

Question Set 4 — Metropolis and Nested Sampling

Question 1

Consider doing basic ‘linear regression’ - fitting of a straight line to data $\{(x_i, y_i)\}_{i=1}^N, N\}$. Make the following assumptions. First, the probability distribution for the data given the parameters:

$$y_i|m, b, \sigma \sim \text{Normal}(mx_i + b, \sigma^2) \quad (1)$$

This is the assumption of ‘gaussian noise’ around the straight line. Astronomers often assume σ is known and different for each data point, perhaps given as a third column in the data file (the ‘error bars’). In this question I’m assuming a constant σ applies to all data points and it is unknown.

Assume the following naive wide priors for the three unknown parameters:

$$m \sim \text{Normal}(0, 1000^2) \quad (2)$$

$$b \sim \text{Normal}(0, 1000^2) \quad (3)$$

$$\ln \sigma \sim \text{Uniform}(-10, 10) \quad (4)$$

Follow these steps to implement this model for the Metropolis algorithm:

1. Copy `transit_model.py` to a new file called `straightline.py`, and work on that.
2. Modify `num_params` and `data` appropriately, and delete un-needed variables.
3. Rewrite `from_prior` and `log_prior`, and modify `jump_sizes`, so they reflect the above prior distributions.
4. Edit `log_likelihood` to be appropriate for this problem.

Run `plain_metropolis.py` and look at trace plots for the parameters. Use the resulting posterior samples to summarise the inferences about m , b , and σ .

For more on fitting straight lines, and the kinds of assumptions that are useful in an astronomy context, see the famous paper by David Hogg.

Question 2

Modify the straight line model so that the noise distribution is potentially heavy tailed:

$$y_i|m, b, \sigma, \nu \sim t(mx_i + b, \sigma^2, \nu) \quad (5)$$

This requires the extra ν parameter. Give it this prior:

$$\ln \nu \sim \text{Uniform}(0, 5). \quad (6)$$

Do your inferences about m , b , and σ change much?

Question 3

Run the two different models in Nested Sampling and get their marginal likelihoods. Interpret the conclusions. You might also want to look at the posterior distribution for ν under Model 2's assumptions.